

ANDREW LAWRENCE

BOSTON, MA

contact@exton.info

<https://exton.info>

EDUCATION	Tufts University BS in Computer Science & Science, Technology, and Society Dean's list, CS Student Council, STS Major Ambassador	Exp. May 2026 Somerville, MA
	La Pontificia Universidad Católica de Chile Tufts in Chile Program	Sept. - Dec. 2024 Santiago, Chile
RESEARCH	Research Assistant	April 2025 - Present
EXPERIENCE	<i>University of Connecticut RIET Lab / Avijit Ghosh</i> <ul style="list-style-type: none">Develop a multi-persona, branching debate web app using SvelteKit to support a study on mitigating online echo chambers, deploying open-source AI models to test the new framework for pluralistic discourse.Design an end-to-end argument mining (NLP) pipeline to process 86K philosophy papers, automating extraction and classification of reasoning patterns to build a novel knowledge graph aimed at improving AI alignment and explainability.	Remote
	Research Assistant <i>University of Michigan School of Information / Ben Green</i> <ul style="list-style-type: none">Critiqued theoretical claims of Algorithmic Realism, applying sociotechnical frameworks to challenge assumptions and directly influence revisions to the work's proposed data science methodology.	Feb. - July. 2024 Remote
INDUSTRY	Software Engineering Intern	Sept. - Dec. 2024
EXPERIENCE	<i>CYSCE S.A. Consultores Empresariales</i> <ul style="list-style-type: none">Modernized a legacy credit-card platform by consolidating customer profiles from 4 business units into one API, directly cutting payment latency.Authored a bilingual (EN/ES) cultural-operations policy playbook adopted by the CEO to steer U.S. market expansion and accelerate cross-division collaboration.	Santiago, Chile
	Digital Transformation Intern <i>Tufts Technology Services</i> <ul style="list-style-type: none">Drove 302% platform engagement growth by leading cross-departmental workflow analyses and deploying tailored SaaS solutions that aligned technical systems with stakeholder goals.Orchestrated a strategic backend migration from legacy systems to AWS (S3, Lambda) after identifying API bottlenecks, cutting infrastructure costs by 15% and boosting system scalability.Developed the organization's first internal analytics suite (Python, Pandas) to translate platform usage data into CIO-level reports, enabling evidence-based decision-making and securing renewed executive investment.	Jan. - Aug. 2024 Somerville, MA
LEADERSHIP	Director (Prev. Technical Fellowship Facilitator)	Jan. 2025 - Present
POSITIONS	<i>Tufts AI Safety</i> <ul style="list-style-type: none">Established Tufts' first AI Safety organization, securing \$20K in funding and scaling the group to 44 members and fellows in its first semester.Designed and led the group's AI policy fellowship, developing an 8-week curriculum for 12 fellows covering topics like compute governance and international coordination. Overseeing our inaugural policy hackathon.Forged partnerships across 2 schools and 4 research labs and hosted a flagship speaker series to integrate responsible AI governance into the university's research and administrative agendas.	Somerville, MA
	Director (Prev. Technical Lead) <i>The Lantern</i> <ul style="list-style-type: none">Co-organize an ongoing AI panel series ("Let's Talk About AI") in partnership with Tufts administration, scaled from an inaugural, 110-attendee event whose findings continue to shape institutional AI policy.Direct a small interdisciplinary team, overseeing and contributing to an annual magazine that provides sociotechnical analyses of technology and science issues (e.g., of surveillance) to 100+ community members.	Jan. 2024 - Present Somerville, MA

AFFILIATIONS	Kairos Pathfinder Program All Tech Is Human MIT AI Alignment	<i>Fellow</i> <i>Member</i> <i>Affiliate</i>
WORK IN PROGRESS	"Contextual Integrity and Differential Privacy; A Sociotechnical Analysis of PETs in AI," <i>Interdis. Honors Thesis</i> . "Mitigating Bias and Echo Chambers in Large Language Models: A Multi-Persona Framework for Nuanced Debates on Sensitive Subjects" "Moral-KG: Building a Knowledge Graph of Philosophical Arguments for Grounded, Transparent AI Reasoning"	
SELECTED WORK	"The Best AI-First Organizations Are Built Bottom-Up" MovieLens100K SVD++ ML Recommender System (RecSys) "TIPS from TwIPS: Integrating Assistive LLM Interp. into Browser Workflows and Evaluating Model Bias" AI Slopsquatting Cybersecurity Policy Memo Logistic Regression UK Lexile Reading-level Classifier "Beyond the Frame of Our Lives: All Light, Everywhere, Surveillance, and the Illusion of Objectivity" "Morality, Capability, and the Ethics of Decision-Making AI" "Pacing Progress; Mitigating Risks in the Rapid Evolution of AI/ML Systems"	<i>July 2025</i> <i>April 2025</i> <i>April 2025</i> <i>March 2025</i> <i>March 2025</i> <i>May 2024</i> <i>Dec. 2023</i> <i>Nov. 2022</i>
EVENTS & TALKS	JumboCode Lightning Talk Series: An AI Safety Perspective on OpenAI's AgentKit RIET Lab Machine Ethics & Reasoning (MERe) Workshop Let's Talk AI: Navigating Academic, Environmental & Societal Impacts	<i>Oct. 2025</i> <i>July 2025</i> <i>April 2025</i>
SKILLS	Theory: Contextual Integrity, Differential Privacy, (Algorithmic) fairness & bias, Argument mining (NLP), LLMs, Fine-tuning (Q/LoRA), AI/ML: LLM, Fine-tuning (Q/LoRA), Transformers, Context engineering, Explainability SWE: Python (Transformers, Scikit-Learn, NumPy, Pandas), C, C++, Git, Docker, REST APIs, AWS Lambda Other: Spanish fluency, Stakeholder management, Workflow analysis, Business Process Model and Notation	